

In [1]:

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
import warnings
import statsmodels.api as sm
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import minmax_scale
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score, accuracy_score, classification_report, precision_recall_
from sklearn import metrics
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.model_selection import GridSearchCV
from category_encoders import TargetEncoder
from sklearn.linear_model import LogisticRegression

```

In [2]:

```
df = pd.read_csv('logistic_regression.csv')
```

EDA

In [3]:

```
df.head()
```

Out[3]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	N
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	N

5 rows × 27 columns

In [4]:

df.shape

Out[4]:

(396030, 27)

In [5]:

df.columns

Out[5]:

```
Index(['loan_amnt', 'term', 'int_rate', 'installment', 'grade', 'sub_grade',
      'emp_title', 'emp_length', 'home_ownership', 'annual_inc',
      'verification_status', 'issue_d', 'loan_status', 'purpose', 'title',
      'dti', 'earliest_cr_line', 'open_acc', 'pub_rec', 'revol_bal',
      'revol_util', 'total_acc', 'initial_list_status', 'application_type',
      'mort_acc', 'pub_rec_bankruptcies', 'address'],
      dtype='object')
```

In [6]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            396030 non-null  float64
1   term                                 396030 non-null  object
2   int_rate                             396030 non-null  float64
3   installment                           396030 non-null  float64
4   grade                                396030 non-null  object
5   sub_grade                            396030 non-null  object
6   emp_title                            373103 non-null  object
7   emp_length                           377729 non-null  object
8   home_ownership                       396030 non-null  object
9   annual_inc                           396030 non-null  float64
10  verification_status                  396030 non-null  object
11  issue_d                              396030 non-null  object
12  loan_status                          396030 non-null  object
13  purpose                              396030 non-null  object
14  title                                394275 non-null  object
15  dti                                  396030 non-null  float64
16  earliest_cr_line                     396030 non-null  object
17  open_acc                             396030 non-null  float64
18  pub_rec                              396030 non-null  float64
19  revol_bal                            396030 non-null  float64
20  revol_util                           395754 non-null  float64
21  total_acc                            396030 non-null  float64
22  initial_list_status                  396030 non-null  object
23  application_type                     396030 non-null  object
24  mort_acc                             358235 non-null  float64
25  pub_rec_bankruptcies                 395495 non-null  float64
26  address                              396030 non-null  object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB
```

In [7]:

```
df.describe(include='all').T
```

Out[7]:

	count	unique	top	freq	mean	std
loan_amnt	396030.0	NaN	NaN	NaN	14113.888089	8357.441341
term	396030	2	36 months	302005	NaN	NaN
int_rate	396030.0	NaN	NaN	NaN	13.6394	4.472157
installment	396030.0	NaN	NaN	NaN	431.849698	250.72779
grade	396030	7	B	116018	NaN	NaN
sub_grade	396030	35	B3	26655	NaN	NaN
emp_title	373103	173105	Teacher	4389	NaN	NaN
emp_length	377729	11	10+ years	126041	NaN	NaN
home_ownership	396030	6	MORTGAGE	198348	NaN	NaN
annual_inc	396030.0	NaN	NaN	NaN	74203.175798	61637.621158
verification_status	396030	3	Verified	139563	NaN	NaN
issue_d	396030	115	Oct-2014	14846	NaN	NaN
loan_status	396030	2	Fully Paid	318357	NaN	NaN
purpose	396030	14	debt_consolidation	234507	NaN	NaN
title	394275	48817	Debt consolidation	152472	NaN	NaN
dti	396030.0	NaN	NaN	NaN	17.379514	18.019092
earliest_cr_line	396030	684	Oct-2000	3017	NaN	NaN
open_acc	396030.0	NaN	NaN	NaN	11.311153	5.137649
pub_rec	396030.0	NaN	NaN	NaN	0.178191	0.530671
revol_bal	396030.0	NaN	NaN	NaN	15844.539853	20591.836109
revol_util	395754.0	NaN	NaN	NaN	53.791749	24.452193
total_acc	396030.0	NaN	NaN	NaN	25.414744	11.886991
initial_list_status	396030	2	f	238066	NaN	NaN
application_type	396030	3	INDIVIDUAL	395319	NaN	NaN
mort_acc	358235.0	NaN	NaN	NaN	1.813991	2.14793
pub_rec_bankruptcies	395495.0	NaN	NaN	NaN	0.121648	0.356174
address	396030	393700	USCGC Smith\r\nFPO AE 70466	8	NaN	NaN

In [59]:

```
#differentiation category and numerical columns
```

In [8]:

```
cat_col = []
con_col = []
for i in df.columns:
    if df[i].nunique() < 40:
        cat_col.append(i)
    else:
        con_col.append(i)
```

In [9]:

```
con_col
```

Out[9]:

```
['loan_amnt',
 'int_rate',
 'installment',
 'emp_title',
 'annual_inc',
 'issue_d',
 'title',
 'dti',
 'earliest_cr_line',
 'open_acc',
 'revol_bal',
 'revol_util',
 'total_acc',
 'address']
```

In [10]:

```
cat_col
```

Out[10]:

```
['term',
 'grade',
 'sub_grade',
 'emp_length',
 'home_ownership',
 'verification_status',
 'loan_status',
 'purpose',
 'pub_rec',
 'initial_list_status',
 'application_type',
 'mort_acc',
 'pub_rec_bankruptcies']
```

In [60]:

```
#finding the missing value columns
```

In [11]:

```
missing_cols = pd.DataFrame(df.isna().sum(),columns=['miss'] )  
missing_cols=missing_cols.loc[missing_cols.miss > 0].index  
missing_cols
```

Out[11]:

```
Index(['emp_title', 'emp_length', 'title', 'revol_util', 'mort_acc',  
      'pub_rec_bankruptcies'],  
      dtype='object')
```

In [61]:

```
#Feature engineering
```

In [12]:

```
df["pub_rec_bankruptcies"]=df["pub_rec_bankruptcies"].apply(lambda x: 1 if x >=1.0 else 0)  
df["pub_rec"]=df["pub_rec"].apply(lambda x: 1 if x >=1.0 else 0)  
df["mort_acc"]=df["mort_acc"].apply(lambda x: 1 if x >=1.0 else 0)  
df['zipcode'] = df['address'].str[-5:]
```

In [13]:

```
df.drop(['address'],inplace=True,axis=1)
```

In [14]:

```
df.nunique()
```

Out[14]:

loan_amnt	1397
term	2
int_rate	566
installment	55706
grade	7
sub_grade	35
emp_title	173105
emp_length	11
home_ownership	6
annual_inc	27197
verification_status	3
issue_d	115
loan_status	2
purpose	14
title	48817
dti	4262
earliest_cr_line	684
open_acc	61
pub_rec	2
revol_bal	55622
revol_util	1226
total_acc	118
initial_list_status	2
application_type	3
mort_acc	2
pub_rec_bankruptcies	2
pincode	10
dtype:	int64

In [15]:

```
cols = df.columns
```

In [16]:

```
df.isna().sum()
```

Out[16]:

```
loan_amnt      0
term           0
int_rate       0
installment    0
grade          0
sub_grade      0
emp_title      22927
emp_length     18301
home_ownership  0
annual_inc     0
verification_status  0
issue_d        0
loan_status    0
purpose        0
title          1755
dti            0
earliest_cr_line  0
open_acc       0
pub_rec        0
revol_bal      0
revol_util     276
total_acc      0
initial_list_status  0
application_type  0
mort_acc       0
pub_rec_bankruptcies  0
pincode        0
dtype: int64
```

In [62]:

```
#checking the quantiles to validate the threshold of quantile
```

In [17]:

```
df.quantile([0.05,0.1,0.2,0.90,0.95,0.97,0.99,1])
```

Out[17]:

	loan_amnt	int_rate	installment	annual_inc	dti	open_acc	pub_rec	revol_bal	rev
0.05	3250.0	6.89	109.5100	28000.0	4.68	5.0	0.0	1685.00	
0.10	5000.0	7.89	158.8600	34000.0	6.90	6.0	0.0	3091.00	
0.20	6625.0	9.67	218.8100	42000.0	10.01	7.0	0.0	5099.80	
0.90	26000.0	19.52	785.4800	120000.0	28.50	18.0	1.0	31470.00	
0.95	30975.0	21.97	925.6000	150000.0	31.58	21.0	1.0	41066.55	
0.97	35000.0	23.28	1028.8013	175000.0	33.34	23.0	1.0	50070.26	
0.99	35000.0	25.28	1202.3730	250000.0	36.43	27.0	1.0	86039.62	
1.00	40000.0	30.99	1533.8100	8706582.0	9999.00	90.0	1.0	1743266.00	

From above quantiles we can see that there is no much outliers in the lower bound and the data is good till 0.97 quantile
We are good to remove the data from 0.97

In [18]:




```
numerical_cols = ['int_rate', 'installment', 'annual_inc', 'dti', 'open_acc', 'pub_rec', 'r
```





In [63]:




```
## univairiate analysis
```


In [19]:

```
for i in numerical_cols:  
    sns.boxplot(x=df[i],y=df['loan_status'])  
    plt.show()
```







In [64]:




```
#removing outliers in the features
```

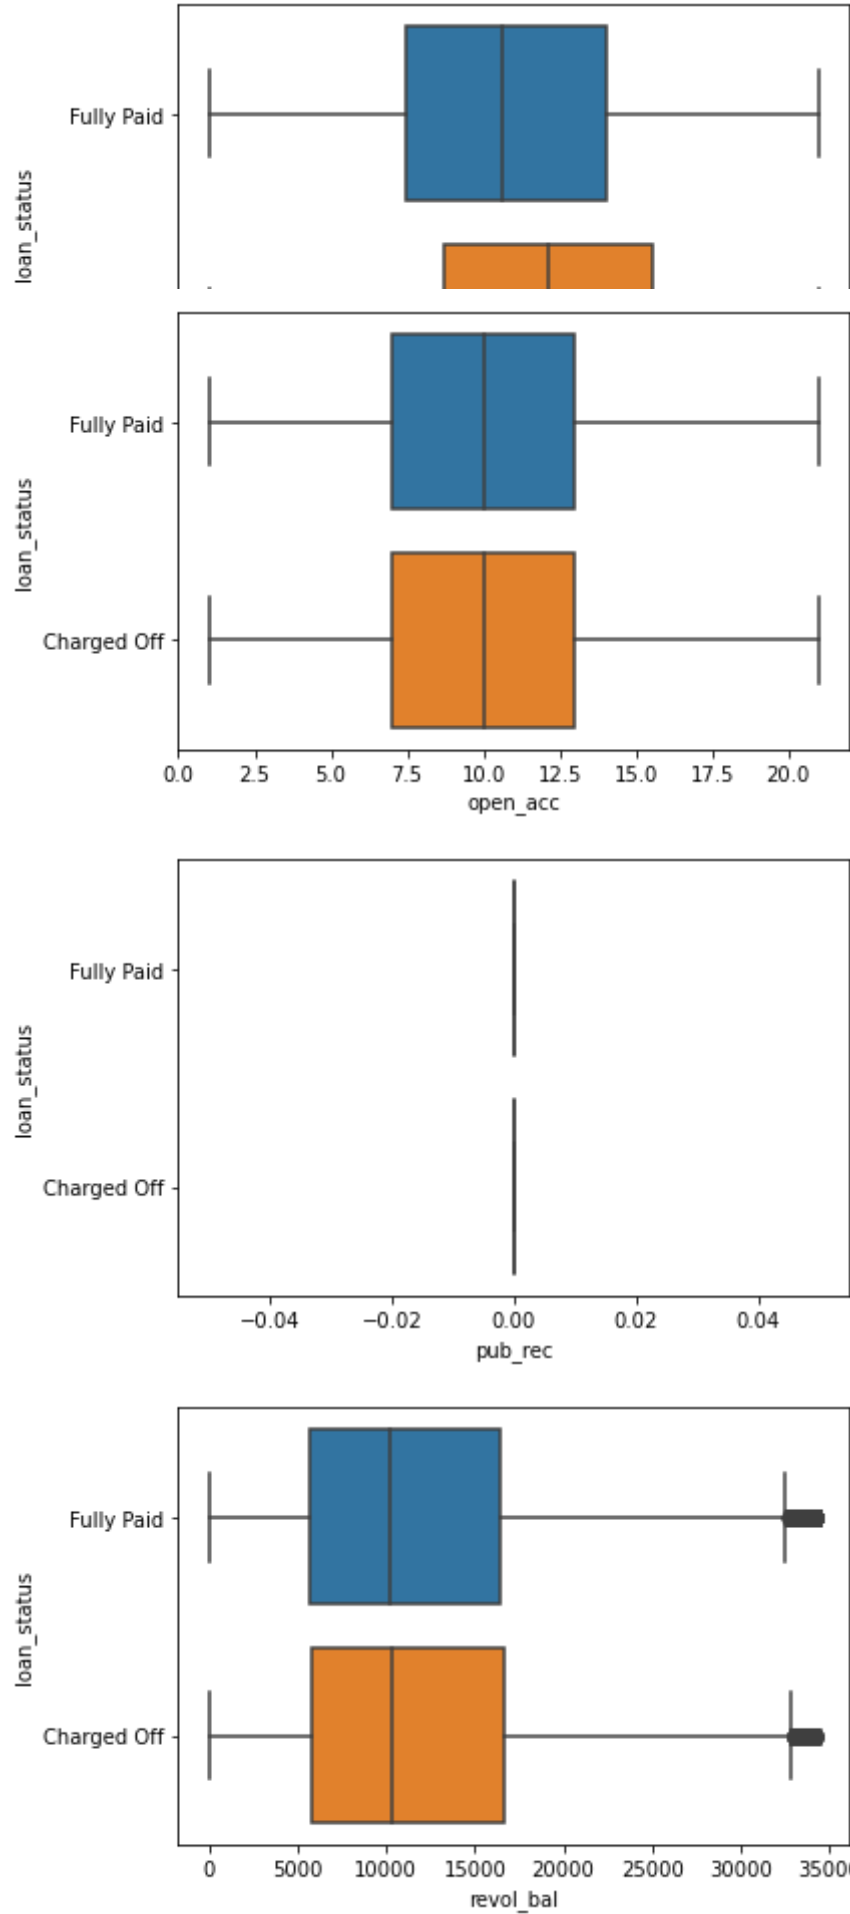
In [20]:




```
df = df[df['annual_inc'] < df['annual_inc'].quantile(0.97)]
df = df[df['int_rate'] < df['int_rate'].quantile(0.97)]
df = df[df['loan_amnt'] < df['loan_amnt'].quantile(0.97)]
df = df[df['installment'] < df['installment'].quantile(0.97)]
df = df[df['annual_inc'] < df['annual_inc'].quantile(0.97)]
df = df[df['open_acc'] < df['open_acc'].quantile(0.97)]
df = df[df['total_acc'] < df['total_acc'].quantile(0.97)]
df = df[df['revol_util'] < df['revol_util'].quantile(0.97)]
df = df[df['dti'] < df['dti'].quantile(0.97)]
df = df[df['pub_rec'] < df['pub_rec'].quantile(0.99)]
df = df[df['revol_bal'] < df['revol_bal'].quantile(0.95)]
```

In [21]:

```
for i in numerical_cols:  
    sns.boxplot(x=df[i],y=df['loan_status'])  
    plt.show()
```







In [65]:


```
#Outliers are treated for the features
```

In [22]:

```
sns.lineplot(x=np.arange(len(df['loan_amnt'][df['loan_status']=='Fully Paid'])),
             y=df['loan_amnt'][df['loan_status']=='Fully Paid'].sort_values()),
sns.lineplot(x=np.arange(len(df['loan_amnt'][df['loan_status']=='Charged Off'])),
             y=df['loan_amnt'][df['loan_status']=='Charged Off'].sort_values())
```

Out[22]:

<AxesSubplot:ylabel='loan_amnt'>



In [66]:

Feature Engineering

In [23]:

df['emp_length'].unique()

Out[23]:

```
array(['4 years', '< 1 year', '6 years', '2 years', '10+ years',
       '7 years', '9 years', '8 years', '5 years', '3 years', '1 year',
       nan], dtype=object)
```

In [24]:

```
# replacing the null values with -1 and greater than 10years with 11 and less than 1 year with 0
df['emp_length'].replace({'< 1 year':'0 years','10+ years':'11 years',np.nan:'-1 years'},inplace=True)
```

In [25]:

```
df['emp_length']=df['emp_length'].str[:5]
df['emp_length']=df['emp_length'].astype('int64')
```

In [26]:

df['emp_title'] = df['emp_title'].replace({np.nan:'Unknown'})

In [27]:

df['title'].replace({np.nan:df['title'].mode()[0]},inplace=True)

In [28]:

```
df['revol_util'].replace({np.nan:df['revol_util'].median()},inplace=True)
```

In [29]:

```
df['grade'].replace({'A':7,'B':6,'C':5,'D':4,'E':3,'F':2,'G':1},inplace=True)
```

In [30]:

```
df['verification_status'].replace({'Verified':1,'Source Verified':1,'Not Verified':0},inplace=True)
```

In [31]:

```
df['issue_month'] = df['issue_d'].str[:3]  
df['issue_year'] = df['issue_d'].str[-4:]  
df.drop(['issue_d'],inplace=True,axis = 1)
```

In [32]:

```
df['earliest_cr_line_month'] = df['earliest_cr_line'].str[:3]  
df['earliest_cr_line_year'] = df['earliest_cr_line'].str[-4:]  
df.drop(['earliest_cr_line'],inplace=True,axis = 1)
```

In [33]:

```
target = df['loan_status']  
target = pd.DataFrame(target,columns=['loan_status'])  
target['loan_status'].replace({'Fully Paid':1,'Charged Off':0},inplace=True)
```

In [34]:

```
df.isna().sum()
```

Out[34]:

loan_amnt	0
term	0
int_rate	0
installment	0
grade	0
sub_grade	0
emp_title	0
emp_length	0
home_ownership	0
annual_inc	0
verification_status	0
loan_status	0
purpose	0
title	0
dti	0
open_acc	0
pub_rec	0
revol_bal	0
revol_util	0
total_acc	0
initial_list_status	0
application_type	0
mort_acc	0
pub_rec_bankruptcies	0
pincode	0
issue_month	0
issue_year	0
earliest_cr_line_month	0
earliest_cr_line_year	0
dtype:	int64

In [67]:

```
# Imputed the features and there are no missing values
```

In [35]:

df.head()

Out[35]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_
1	8000.0	36 months	11.99	265.68	6	B5	Credit analyst	4	N
2	15600.0	36 months	10.49	506.97	6	B3	Statistician	0	
3	7200.0	36 months	6.49	220.65	7	A2	Client Advocate	6	
6	18000.0	36 months	5.32	542.07	7	A1	Software Development Engineer	2	N
7	13000.0	36 months	11.14	426.47	6	B2	Office Depot	11	

5 rows × 29 columns

In [36]:

df['home_ownership'].value_counts()

Out[36]:

```

MORTGAGE    109345
RENT         106567
OWN          22535
OTHER         91
NONE         21
ANY           2
Name: home_ownership, dtype: int64

```

The value count for Other, None, Any is very low compared to others, lets combine them.

In [37]:

df['home_ownership'] = df['home_ownership'].replace({'NONE': 'OTHER', 'ANY': 'OTHER'})

In [38]:

df['home_ownership'].value_counts()

Out[38]:

```

MORTGAGE    109345
RENT         106567
OWN          22535
OTHER         114
Name: home_ownership, dtype: int64

```

In [39]:


```
## To reduce data inconsistency
df['title'] = df['title'].str.lower()
df['purpose'] = df['purpose'].str.lower()
df['emp_title'] = df['emp_title'].str.lower()
```

In [40]:

```
plt.xticks(rotation =45)
plt.plot(df['emp_title'].value_counts()[:10])
```

Out[40]:

[<matplotlib.lines.Line2D at 0x2a387c85780>]



Teacher and manager are the most common professions,


In [41]:

```
sns.countplot(df['loan_status'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning

Out[41]:

```
<AxesSubplot:xlabel='loan_status', ylabel='count'>
```



In [42]:

```
x = df['loan_status'][df['loan_status']=='Fully Paid'].value_counts()
y = df['loan_status'][df['loan_status']=='Charged Off'].value_counts()

print('Percentage of fully paid customers is ' + str(np.round(list(x)[0]/(list(y)[0]+list(x)
```

Percentage of fully paid customers is 81.84%

In [57]:

```
x = df['grade'][(df['grade']==7) & (df['loan_status'] == 'Fully Paid')].value_counts()
y = df['grade'][(df['grade']==7) & (df['loan_status'] == 'Charged Off')].value_counts()
print('Percentage of fully paid customers is ' + str(np.round(list(x)[0]/(list(y)[0]+list(x)
```


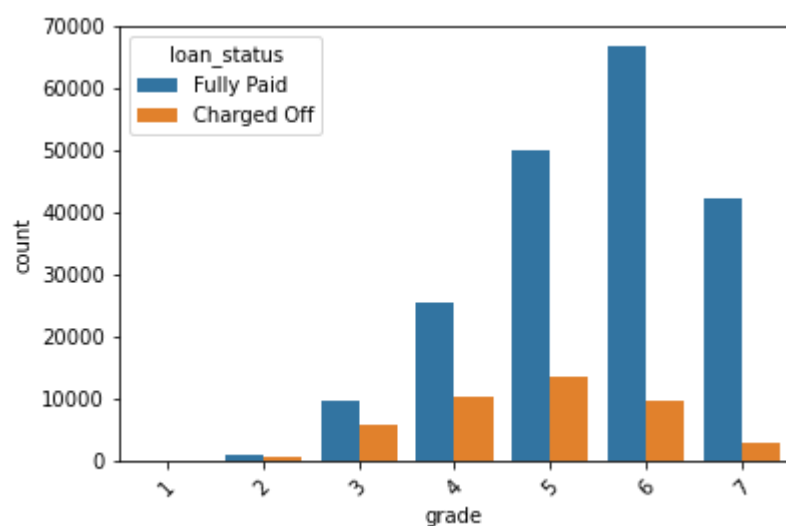
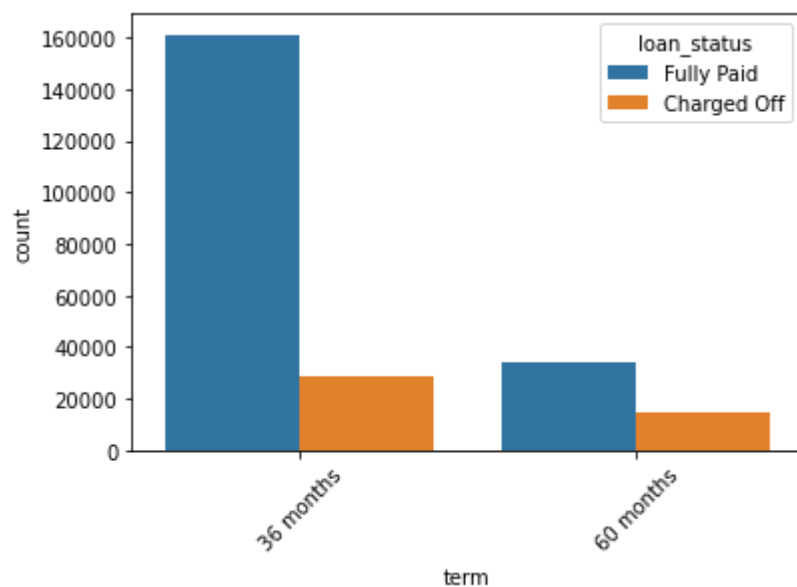
Percentage of fully paid customers is 93.71%

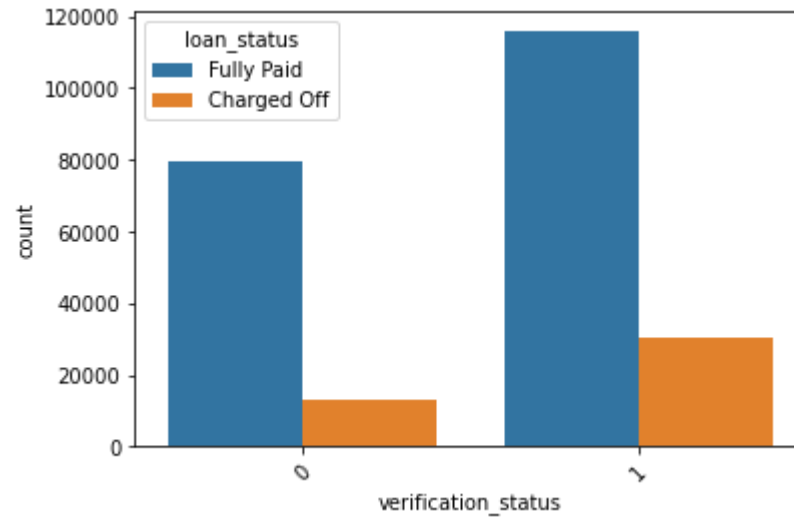
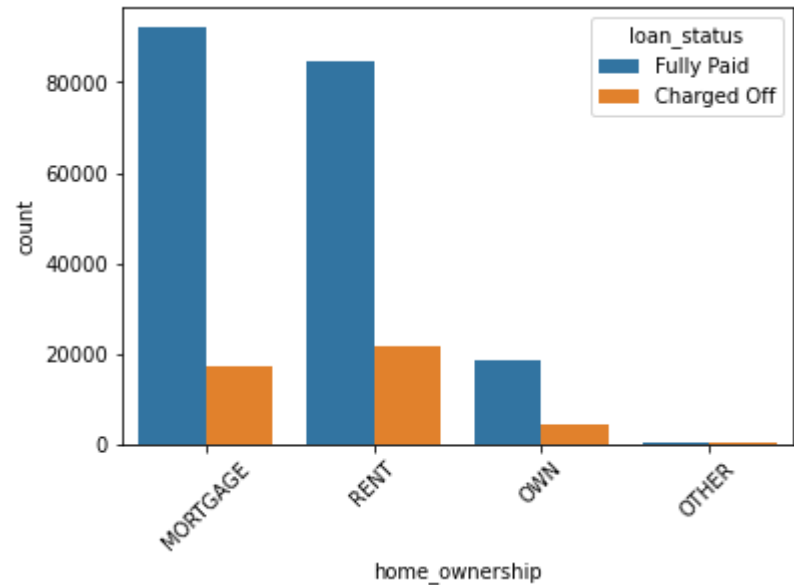

In [68]:

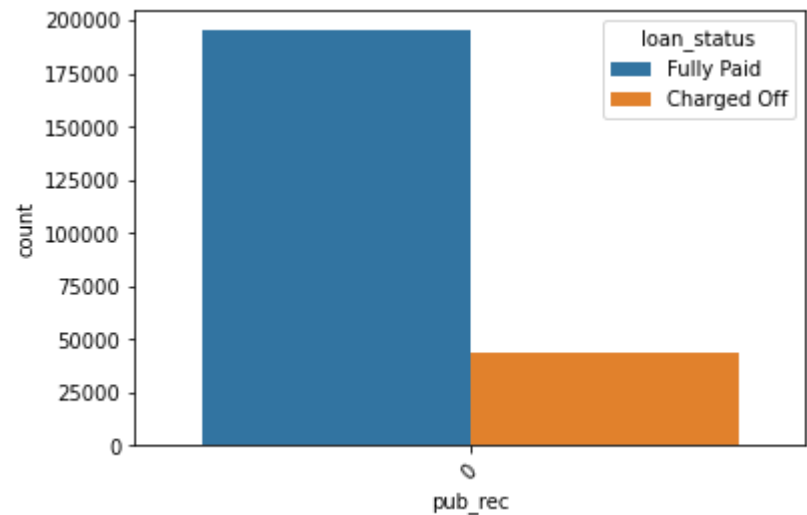
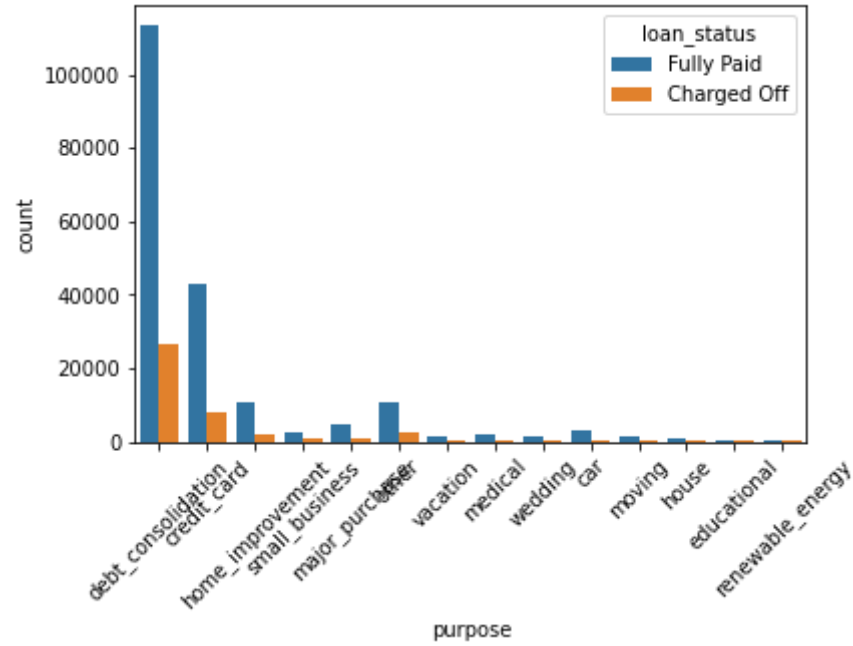

```
## BiVariate Analysis
```


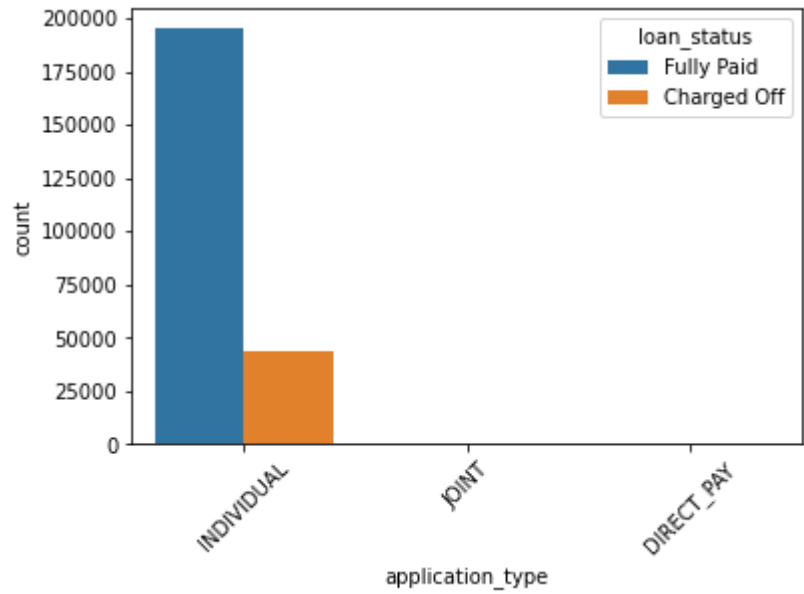


In [43]:

```
for i in cat_col:  
    sns.countplot(x=df[i],hue = df['loan_status'])  
    plt.xticks(rotation =45)  
    plt.show()
```











Most of the customers opt for 36 months term
Percentage of customers who Fully paid the loan is greater in 36months term than compared to 60 months
Customers with Grade E mostly dont pay the loan completely
Customers with Grade B and A pays off the loan
Customers with emp_length having 10+ years takes more amount of loan
Home_ownership with Mortgage takes more number of loans and payoff
Most of the customers takes loans as Individual compared to joint and direct_pay

In [44]:

```
con_col
```

Out[44]:

```
['loan_amnt',  
 'int_rate',  
 'installment',  
 'emp_title',  
 'annual_inc',  
 'issue_d',  
 'title',  
 'dti',  
 'earliest_cr_line',  
 'open_acc',  
 'revol_bal',  
 'revol_util',  
 'total_acc',  
 'address']
```

In [45]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 238561 entries, 1 to 396029
Data columns (total 29 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            238561 non-null float64
1   term                                 238561 non-null object
2   int_rate                             238561 non-null float64
3   installment                          238561 non-null float64
4   grade                                238561 non-null int64
5   sub_grade                            238561 non-null object
6   emp_title                            238561 non-null object
7   emp_length                           238561 non-null int64
8   home_ownership                       238561 non-null object
9   annual_inc                           238561 non-null float64
10  verification_status                  238561 non-null int64
11  loan_status                           238561 non-null object
12  purpose                               238561 non-null object
13  title                                238561 non-null object
14  dti                                   238561 non-null float64
15  open_acc                             238561 non-null float64
16  pub_rec                              238561 non-null int64
17  revol_bal                            238561 non-null float64
18  revol_util                           238561 non-null float64
19  total_acc                            238561 non-null float64
20  initial_list_status                   238561 non-null object
21  application_type                      238561 non-null object
22  mort_acc                              238561 non-null int64
23  pub_rec_bankruptcies                  238561 non-null int64
24  pincode                               238561 non-null object
25  issue_month                           238561 non-null object
26  issue_year                            238561 non-null object
27  earliest_cr_line_month                 238561 non-null object
28  earliest_cr_line_year                  238561 non-null object
dtypes: float64(9), int64(6), object(14)
memory usage: 62.7+ MB
```

In [46]:

```
df[['int_rate', 'installment', 'annual_inc', 'dti', 'open_acc', 'pub_rec', 'revol_bal', 're
```

In [47]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 238561 entries, 1 to 396029
Data columns (total 29 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   loan_amnt             238561 non-null float64
 1   term                  238561 non-null object
 2   int_rate              238561 non-null float64
 3   installment           238561 non-null float64
 4   grade                 238561 non-null int64
 5   sub_grade             238561 non-null object
 6   emp_title             238561 non-null object
 7   emp_length           238561 non-null int64
 8   home_ownership        238561 non-null object
 9   annual_inc            238561 non-null float64
10  verification_status    238561 non-null int64
11  loan_status           238561 non-null object
12  purpose               238561 non-null object
13  title                 238561 non-null object
14  dti                   238561 non-null float64
15  open_acc              238561 non-null float64
16  pub_rec               238561 non-null float64
17  revol_bal             238561 non-null float64
18  revol_util            238561 non-null float64
19  total_acc             238561 non-null float64
20  initial_list_status    238561 non-null object
21  application_type       238561 non-null object
22  mort_acc              238561 non-null float64
23  pub_rec_bankruptcies  238561 non-null int64
24  pincode               238561 non-null object
25  issue_month           238561 non-null object
26  issue_year            238561 non-null object
27  earliest_cr_line_month 238561 non-null object
28  earliest_cr_line_year  238561 non-null object
dtypes: float64(11), int64(4), object(14)
memory usage: 62.7+ MB
```

In [48]:

df['verification_status'].unique()

Out[48]:

array([0, 1], dtype=int64)

In [49]:

```
df[['int_rate',
    'installment',
    'annual_inc',
    'dti',
    'open_acc',
    'pub_rec',
    'revol_bal',
    'revol_util',
    'total_acc',
    'mort_acc']].head()
```

Out[49]:

	int_rate	installment	annual_inc	dti	open_acc	pub_rec	revol_bal	revol_util	total_acc
1	11.99	265.68	65000.0	22.05	17.0	0.0	20131.0	53.3	27.0
2	10.49	506.97	43057.0	12.79	13.0	0.0	11987.0	92.2	26.0
3	6.49	220.65	54000.0	2.60	6.0	0.0	5472.0	21.5	13.0
6	5.32	542.07	125000.0	1.36	8.0	0.0	4178.0	4.9	25.0
7	11.14	426.47	46000.0	26.87	11.0	0.0	13425.0	64.5	15.0

In [101]:

df.head()

Out[101]:

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_
1	8000.0	36 months	11.99	265.68	6	B5	credit analyst	4	N
2	15600.0	36 months	10.49	506.97	6	B3	statistician	0	
3	7200.0	36 months	6.49	220.65	7	A2	client advocate	6	
6	18000.0	36 months	5.32	542.07	7	A1	software development engineer	2	N
7	13000.0	36 months	11.14	426.47	6	B2	office depot	11	

5 rows × 29 columns

In [102]:

df['purpose'].unique()

Out[102]:

```
array(['debt_consolidation', 'credit_card', 'home_improvement',
    'small_business', 'major_purchase', 'other', 'vacation', 'medical',
    'wedding', 'car', 'moving', 'house', 'educational',
    'renewable_energy'], dtype=object)
```

In [103]:

```
df.nunique()
```

Out[103]:

loan_amnt	1270
term	2
int_rate	476
installment	40809
grade	7
sub_grade	35
emp_title	106288
emp_length	12
home_ownership	4
annual_inc	17504
verification_status	2
loan_status	2
purpose	14
title	31070
dti	3317
open_acc	21
pub_rec	1
revol_bal	32808
revol_util	1002
total_acc	45
initial_list_status	2
application_type	3
mort_acc	2
pub_rec_bankruptcies	1
pincode	10
issue_month	12
issue_year	10
earliest_cr_line_month	12
earliest_cr_line_year	63
dtype:	int64

In [69]:

```
# One hot Encoding
```

In [104]:

```
one_hot_cols=['term','verification_status','pub_rec','initial_list_status','application_type']
```

In [105]:

```

cont = pd.get_dummies(df['term'],prefix='term',drop_first=True)
#Adding the results to the master dataframe
df = pd.concat([df,cont],axis=1)

cont = pd.get_dummies(df['verification_status'],prefix='verification_status',drop_first=True)
#Adding the results to the master dataframe
df = pd.concat([df,cont],axis=1)

cont = pd.get_dummies(df['pub_rec'],prefix='pub_rec',drop_first=True)
#Adding the results to the master dataframe
df = pd.concat([df,cont],axis=1)

cont = pd.get_dummies(df['initial_list_status'],prefix='initial_list_status',drop_first=True)
#Adding the results to the master dataframe
df = pd.concat([df,cont],axis=1)

cont = pd.get_dummies(df['application_type'],prefix='application_type',drop_first=True)
#Adding the results to the master dataframe
df = pd.concat([df,cont],axis=1)

cont = pd.get_dummies(df['mort_acc'],prefix='mort_acc',drop_first=True)
#Adding the results to the master dataframe
df = pd.concat([df,cont],axis=1)

cont = pd.get_dummies(df['pub_rec_bankruptcies'],prefix='pub_rec_bankruptcies',drop_first=True)
#Adding the results to the master dataframe
df = pd.concat([df,cont],axis=1)

```

In [106]:

```
df.drop(one_hot_cols,inplace=True,axis = 1)
```


In [107]:

```
df.head()
```

Out[107]:

	loan_amnt	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownersh
1	8000.0	11.99	265.68	6	B5	credit analyst	4	MORTGAG
2	15600.0	10.49	506.97	6	B3	statistician	0	REN
3	7200.0	6.49	220.65	7	A2	client advocate	6	REN
6	18000.0	5.32	542.07	7	A1	software development engineer	2	MORTGAG
7	13000.0	11.14	426.47	6	B2	office depot	11	REN

5 rows × 28 columns



In [108]:

```
df['loan_status'].replace({'Fully Paid':0, 'Charged Off':1}, inplace=True)
```

In [109]:

```
df['loan_status']
```

Out[109]:

```
1      0
2      0
3      0
6      0
7      0
..
396023  0
396025  0
396027  0
396028  0
396029  0
Name: loan_status, Length: 238561, dtype: int64
```

In [70]:

```
# Target Encoding
```

In [110]:

```
te = TargetEncoder(return_df=False)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\category_encoders\target_encoder.py:124: FutureWarning: Default parameter min_samples_leaf will change in version 2.6. See https://github.com/scikit-learn-contrib/category\_encoders/issues/327 (https://github.com/scikit-learn-contrib/category\_encoders/issues/327)
category=FutureWarning)
C:\ProgramData\Anaconda3\lib\site-packages\category_encoders\target_encoder.py:129: FutureWarning: Default parameter smoothing will change in version 2.6. See https://github.com/scikit-learn-contrib/category\_encoders/issues/327 (https://github.com/scikit-learn-contrib/category\_encoders/issues/327)
category=FutureWarning)
```

In [111]:

```
df.nunique()
```

Out[111]:

loan_amnt	1270
int_rate	476
installment	40809
grade	7
sub_grade	35
emp_title	106288
emp_length	12
home_ownership	4
annual_inc	17504
loan_status	2
purpose	14
title	31070
dti	3317
open_acc	21
revol_bal	32808
revol_util	1002
total_acc	45
pincode	10
issue_month	12
issue_year	10
earliest_cr_line_month	12
earliest_cr_line_year	63
term_ 60 months	2
verification_status_1	2
initial_list_status_w	2
application_type_INDIVIDUAL	2
application_type_JOINT	2
mort_acc_1.0	2
dtype:	int64

In [112]:

```
te_cols = ['grade', 'sub_grade', 'title', 'emp_title', 'pincode', 'emp_length', 'home_ownership',
for col in te_cols:
    df[col]= df[col].astype('category')
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 238561 entries, 1 to 396029
Data columns (total 28 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   loan_amnt                            238561 non-null float64
1   int_rate                            238561 non-null float64
2   installment                          238561 non-null float64
3   grade                               238561 non-null category
4   sub_grade                           238561 non-null category
5   emp_title                           238561 non-null category
6   emp_length                          238561 non-null category
7   home_ownership                      238561 non-null category
8   annual_inc                          238561 non-null float64
9   loan_status                         238561 non-null int64
10  purpose                             238561 non-null category
11  title                               238561 non-null category
12  dti                                 238561 non-null float64
13  open_acc                           238561 non-null float64
14  revol_bal                           238561 non-null float64
15  revol_util                           238561 non-null float64
16  total_acc                           238561 non-null float64
17  pincode                             238561 non-null category
18  issue_month                         238561 non-null category
19  issue_year                          238561 non-null category
20  earliest_cr_line_month              238561 non-null category
21  earliest_cr_line_year               238561 non-null category
22  term_ 60 months                     238561 non-null uint8
23  verification_status_1               238561 non-null uint8
24  initial_list_status_w               238561 non-null uint8
25  application_type_INDIVIDUAL         238561 non-null uint8
26  application_type_JOINT              238561 non-null uint8
27  mort_acc_1.0                        238561 non-null uint8
dtypes: category(12), float64(9), int64(1), uint8(6)
memory usage: 39.2 MB
```

In [113]:

```
for col in te_cols:
    te=TargetEncoder()
    te.fit(X=df[col],y=df['loan_status'])
    values = te.transform(df[col],override_return_df=True)
    df = pd.concat([df,values],axis=1)
```

In [114]:

```
df = df.T[~df.T.index.duplicated(keep='last')].T
```

In [115]:

```
df.head()
```

Out[115]:

	loan_amnt	int_rate	installment	annual_inc	loan_status	dti	open_acc	revol_bal	revol_util
1	8000.0	11.99	265.68	65000.0	0	22.05	17.0	20131.0	5
2	15600.0	10.49	506.97	43057.0	0	12.79	13.0	11987.0	9
3	7200.0	6.49	220.65	54000.0	0	2.6	6.0	5472.0	2
6	18000.0	5.32	542.07	125000.0	0	1.36	8.0	4178.0	
7	13000.0	11.14	426.47	46000.0	0	26.87	11.0	13425.0	6

5 rows × 28 columns

In [116]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 238561 entries, 1 to 396029
Data columns (total 28 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   loan_amnt                            238561 non-null  object
 1   int_rate                             238561 non-null  object
 2   installment                           238561 non-null  object
 3   annual_inc                           238561 non-null  object
 4   loan_status                           238561 non-null  object
 5   dti                                   238561 non-null  object
 6   open_acc                             238561 non-null  object
 7   revol_bal                             238561 non-null  object
 8   revol_util                             238561 non-null  object
 9   total_acc                             238561 non-null  object
10  term_60_months                         238561 non-null  object
11  verification_status_1                 238561 non-null  object
12  initial_list_status_w                 238561 non-null  object
13  application_type_INDIVIDUAL           238561 non-null  object
14  application_type_JOINT                 238561 non-null  object
15  mort_acc_1.0                           238561 non-null  object
16  grade                                 238561 non-null  object
17  sub_grade                             238561 non-null  object
18  title                                 238561 non-null  object
19  emp_title                             238561 non-null  object
20  pincode                               238561 non-null  object
21  emp_length                             238561 non-null  object
22  home_ownership                         238561 non-null  object
23  purpose                               238561 non-null  object
24  issue_month                           238561 non-null  object
25  issue_year                             238561 non-null  object
26  earliest_cr_line_month                 238561 non-null  object
27  earliest_cr_line_year                 238561 non-null  object
dtypes: object(28)
memory usage: 60.8+ MB
```

In [117]:

```
target = df['loan_status']
target = pd.DataFrame(target, columns=['loan_status'])
target['loan_status'].replace({'Fully Paid':1, 'Charged Off':0}, inplace=True)
```

In [118]:

```
# df.drop(['loan_status'], inplace = True, axis =1)
```

In [119]:

```
#Normalisation
for col in df.columns:
    df[col] = minmax_scale(df[col])
df.head()
```

Out[119]:

	loan_amnt	int_rate	installment	annual_inc	loan_status	dti	open_acc	revol_bal	re
1	0.221239	0.371795	0.290129	0.451865	0.0	0.664557	0.80	0.585084	0
2	0.445428	0.288183	0.570790	0.289320	0.0	0.385473	0.60	0.348388	0
3	0.197640	0.065217	0.237752	0.370381	0.0	0.078360	0.25	0.159037	0
6	0.516224	0.000000	0.611618	0.896323	0.0	0.040989	0.35	0.121429	0
7	0.368732	0.324415	0.477155	0.311120	0.0	0.809825	0.50	0.390182	0

5 rows × 28 columns

In [120]:

```
df.columns
```

Out[120]:

```
Index(['loan_amnt', 'int_rate', 'installment', 'annual_inc', 'loan_status',
      'dti', 'open_acc', 'revol_bal', 'revol_util', 'total_acc',
      'term_ 60 months', 'verification_status_1', 'initial_list_status_w',
      'application_type_INDIVIDUAL', 'application_type_JOINT', 'mort_acc_1.
0',
      'grade', 'sub_grade', 'title', 'emp_title', 'pincode', 'emp_length',
      'home_ownership', 'purpose', 'issue_month', 'issue_year',
      'earliest_cr_line_month', 'earliest_cr_line_year'],
      dtype='object')
```

In [71]:


```
#Heatmap
```

In [121]:

```
plt.figure(figsize=(30,10))
sns.heatmap(df.corr(),annot=True)
```

Out[121]:

<AxesSubplot:>



Correlation for the columns loan_amnt,subgrade,open_acnt,mort_acc,revol_bal is high, hence removing the columns.

In [122]:

df.columns

Out[122]:

```
Index(['loan_amnt', 'int_rate', 'installment', 'annual_inc', 'loan_status',
       'dti', 'open_acc', 'revol_bal', 'revol_util', 'total_acc',
       'term_60_months', 'verification_status_1', 'initial_list_status_w',
       'application_type_INDIVIDUAL', 'application_type_JOINT', 'mort_acc_1.
0',
       'grade', 'sub_grade', 'title', 'emp_title', 'pincode', 'emp_length',
       'home_ownership', 'purpose', 'issue_month', 'issue_year',
       'earliest_cr_line_month', 'earliest_cr_line_year'],
      dtype='object')
```

In [123]:

```
df.drop(['loan_amnt', 'sub_grade', 'open_acc', 'mort_acc_1.0', 'revol_bal'],axis=1, inplace = T
```

Model Building

In [124]:

```

X = df.drop(['loan_status'],axis=1)
y = df['loan_status']

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob=logsk.predict_proba(X_test)

#Probabilities
print(y_pred,y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test,y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[9.32155336e-01 6.78446643e-02]
 [8.45370448e-01 1.54629552e-01]
 [9.99554381e-01 4.45618759e-04]
 ...
 [9.98900046e-01 1.09995421e-03]
 [8.92975114e-01 1.07024886e-01]
 [8.41091096e-01 1.58908904e-01]]

```

```

****Score is :0.9290895841770686****

```

	precision	recall	f1-score	support
0.0	0.90	0.99	0.94	58474
1.0	0.90	0.53	0.67	13095
accuracy			0.90	71569
macro avg	0.90	0.76	0.81	71569
weighted avg	0.90	0.90	0.89	71569

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.p
y:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,

In [125]:

```
df.shape
```

Out[125]:

(238561, 23)

In [126]:

```
col = ['int_rate', 'installment', 'annual_inc', 'dti', 'revol_util',  
       'total_acc', 'term_ 60 months', 'verification_status_1',  
       'initial_list_status_w', 'application_type_INDIVIDUAL',  
       'application_type_JOINT', 'grade', 'title', 'emp_title', 'pincode',  
       'emp_length', 'home_ownership', 'purpose', 'issue_month', 'issue_year',  
       'earliest_cr_line_month', 'earliest_cr_line_year']
```


In [127]:

```
#Logistic regression model
logm2 = sm.GLM(y_train,(sm.add_constant(X_train)),family=sm.families.Binomial())
logm2.fit().summary()
```

Out[127]:

Generalized Linear Model Regression Results

Dep. Variable:	loan_status	No. Observations:	166992
Model:	GLM	Df Residuals:	166969
Model Family:	Binomial	Df Model:	22
Link Function:	logit	Scale:	1.0000
Method:	IRLS	Log-Likelihood:	-37369.
Date:	Mon, 28 Nov 2022	Deviance:	74739.
Time:	22:29:32	Pearson chi2:	8.93e+04
No. Iterations:	12	Pseudo R-squ. (CS):	0.3925
Covariance Type:	nonrobust		

	coef	std err	z	P> z	[0.025	0.975]
const	-12.6853	1.206	-10.518	0.000	-15.049	-10.322
int_rate	1.4341	0.128	11.235	0.000	1.184	1.684
installment	0.3605	0.049	7.305	0.000	0.264	0.457
annual_inc	-0.7617	0.064	-11.821	0.000	-0.888	-0.635
dti	0.5570	0.046	12.132	0.000	0.467	0.647
revol_util	0.4152	0.042	9.902	0.000	0.333	0.497
total_acc	-0.0561	0.053	-1.068	0.286	-0.159	0.047
term_ 60 months	0.5589	0.024	23.180	0.000	0.512	0.606
verification_status_1	0.0954	0.021	4.449	0.000	0.053	0.137
initial_list_status_w	-0.1496	0.020	-7.424	0.000	-0.189	-0.110
application_type_INDIVIDUAL	-0.4616	1.171	-0.394	0.693	-2.756	1.833
application_type_JOINT	-0.6873	1.214	-0.566	0.571	-3.067	1.693
grade	0.3686	0.099	3.730	0.000	0.175	0.562
title	6.3522	0.193	32.941	0.000	5.974	6.730
emp_title	6.8459	0.098	70.095	0.000	6.654	7.037
pincode	38.3579	1.448	26.494	0.000	35.520	41.196
emp_length	0.2301	0.041	5.646	0.000	0.150	0.310
home_ownership	0.2299	0.021	10.895	0.000	0.189	0.271
purpose	-0.2243	0.077	-2.914	0.004	-0.375	-0.073
issue_month	0.1423	0.032	4.384	0.000	0.079	0.206
issue_year	0.4460	0.030	15.106	0.000	0.388	0.504
earliest_cr_line_month	0.0601	0.032	1.874	0.061	-0.003	0.123
earliest_cr_line_year	0.7592	0.378	2.006	0.045	0.018	1.501

In [128]:

```

xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_v
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data

```

Out[128]:

	feature	VIF
9	application_type_INDIVIDUAL	157.18
21	earliest_cr_line_year	95.82
0	int_rate	42.76
11	grade	23.96
12	title	14.96
17	purpose	13.15
2	annual_inc	9.17
5	total_acc	8.19
3	dti	7.60
4	revol_util	7.38
1	installment	6.01
18	issue_month	5.34
13	emp_title	4.89
19	issue_year	4.68
20	earliest_cr_line_month	3.55
7	verification_status_1	2.98
16	home_ownership	2.48
8	initial_list_status_w	1.77
6	term_ 60 months	1.70
14	pincode	1.63
15	emp_length	1.32
10	application_type_JOINT	1.11

In [129]:

```
X_train.columns
```

Out[129]:

```
Index(['int_rate', 'installment', 'annual_inc', 'dti', 'revol_util',  
      'total_acc', 'term_ 60 months', 'verification_status_1',  
      'initial_list_status_w', 'application_type_INDIVIDUAL',  
      'application_type_JOINT', 'grade', 'title', 'emp_title', 'pincode',  
      'emp_length', 'home_ownership', 'purpose', 'issue_month', 'issue_yea  
r',  
      'earliest_cr_line_month', 'earliest_cr_line_year'],  
      dtype='object')
```

In [130]:

```

## P-value and VIF for application_type_INDIVIDUAL is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL'], axis=1)
y = df['loan_status']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob = logsk.predict_proba(X_test)

#Probabilities
print(y_pred, y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test, y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[9.31572429e-01 6.84275706e-02]
 [8.43141499e-01 1.56858501e-01]
 [9.99569204e-01 4.30796355e-04]
 ...
 [9.98949691e-01 1.05030924e-03]
 [8.93197713e-01 1.06802287e-01]
 [8.39207930e-01 1.60792070e-01]]

```

```

****Score is :0.9290930306199405****

```

	precision	recall	f1-score	support
0.0	0.90	0.99	0.94	58474
1.0	0.90	0.53	0.67	13095
accuracy			0.90	71569
macro avg	0.90	0.76	0.80	71569
weighted avg	0.90	0.90	0.89	71569

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_logistic.p
y:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,

In [131]:

```
#Logistic regression model
logm3 = sm.GLM(y_train,(sm.add_constant(X_train)),family=sm.families.Binomial())
logm3.fit().summary()
```

Out[131]:

Generalized Linear Model Regression Results

Dep. Variable:	loan_status		No. Observations:	166992			
Model:	GLM		Df Residuals:	166970			
Model Family:	Binomial		Df Model:	21			
Link Function:	logit		Scale:	1.0000			
Method:	IRLS		Log-Likelihood:	-37369.			
Date:	Mon, 28 Nov 2022		Deviance:	74739.			
Time:	22:29:42		Pearson chi2:	8.93e+04			
No. Iterations:	12		Pseudo R-squ. (CS):	0.3925			
Covariance Type:	nonrobust						
		coef	std err	z	P> z	[0.025	0.975]
	const	-13.1470	0.290	-45.366	0.000	-13.715	-12.579
	int_rate	1.4340	0.128	11.234	0.000	1.184	1.684
	installment	0.3604	0.049	7.303	0.000	0.264	0.457
	annual_inc	-0.7614	0.064	-11.818	0.000	-0.888	-0.635
	dti	0.5573	0.046	12.139	0.000	0.467	0.647
	revol_util	0.4151	0.042	9.900	0.000	0.333	0.497
	total_acc	-0.0561	0.053	-1.068	0.286	-0.159	0.047
	term_60 months	0.5588	0.024	23.178	0.000	0.512	0.606
	verification_status_1	0.0954	0.021	4.449	0.000	0.053	0.137
	initial_list_status_w	-0.1496	0.020	-7.424	0.000	-0.189	-0.110
application_type_JOINT		-0.2258	0.324	-0.697	0.486	-0.861	0.409
	grade	0.3689	0.099	3.733	0.000	0.175	0.563
	title	6.3523	0.193	32.942	0.000	5.974	6.730
	emp_title	6.8458	0.098	70.094	0.000	6.654	7.037
	pincode	38.3576	1.448	26.494	0.000	35.520	41.195
	emp_length	0.2301	0.041	5.646	0.000	0.150	0.310
	home_ownership	0.2299	0.021	10.895	0.000	0.189	0.271
	purpose	-0.2244	0.077	-2.916	0.004	-0.375	-0.074
	issue_month	0.1424	0.032	4.385	0.000	0.079	0.206
	issue_year	0.4459	0.030	15.102	0.000	0.388	0.504
earliest_cr_line_month		0.0601	0.032	1.875	0.061	-0.003	0.123
earliest_cr_line_year		0.7592	0.378	2.006	0.045	0.018	1.501

In [132]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_v
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[132]:

	feature	VIF
0	int_rate	41.93
20	earliest_cr_line_year	32.19
10	grade	23.26
11	title	14.55
16	purpose	12.60
2	annual_inc	8.69
3	dti	7.55
5	total_acc	7.31
4	revol_util	7.15
1	installment	6.00
17	issue_month	5.21
12	emp_title	4.80
18	issue_year	4.68
19	earliest_cr_line_month	3.52
7	verification_status_1	2.98
15	home_ownership	2.47
8	initial_list_status_w	1.77
6	term_60 months	1.69
13	pincode	1.63
14	emp_length	1.31
9	application_type_JOINT	1.00

In [133]:

```

## P-value and VIF for earliest_cr_line_year is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL', 'earliest_cr_line_year'], axis=1)
y = df['loan_status']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob = logsk.predict_proba(X_test)

#Probabilities
print(y_pred, y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test, y_pred))

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model_logistic.p
y:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,

[0. 0. 0. ... 0. 0. 0.] [[9.31389072e-01 6.86109279e-02]
 [8.41884440e-01 1.58115560e-01]
 [9.99572733e-01 4.27267242e-04]
 ...
 [9.98964834e-01 1.03516631e-03]
 [8.93529615e-01 1.06470385e-01]
 [8.37612545e-01 1.62387455e-01]]
****Score is :0.9290824457698166****

```

	precision	recall	f1-score	support
0.0	0.90	0.99	0.94	58474
1.0	0.90	0.53	0.67	13095
accuracy			0.90	71569
macro avg	0.90	0.76	0.81	71569
weighted avg	0.90	0.90	0.89	71569

In [134]:

```
#Logistic regression model
logm4 = sm.GLM(y_train,(sm.add_constant(X_train)),family=sm.families.Binomial())
logm4.fit().summary()
```

Out[134]:

Generalized Linear Model Regression Results

Dep. Variable:	loan_status	No. Observations:	166992			
Model:	GLM	Df Residuals:	166971			
Model Family:	Binomial	Df Model:	20			
Link Function:	logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-37371.			
Date:	Mon, 28 Nov 2022	Deviance:	74743.			
Time:	22:29:52	Pearson chi2:	8.93e+04			
No. Iterations:	12	Pseudo R-squ. (CS):	0.3925			
Covariance Type:	nonrobust					
	coef	std err	z	P> z	[0.025	0.975]
const	-12.9638	0.275	-47.146	0.000	-13.503	-12.425
int_rate	1.4395	0.128	11.281	0.000	1.189	1.690
installment	0.3564	0.049	7.230	0.000	0.260	0.453
annual_inc	-0.7675	0.064	-11.925	0.000	-0.894	-0.641
dti	0.5598	0.046	12.200	0.000	0.470	0.650
revol_util	0.4066	0.042	9.750	0.000	0.325	0.488
total_acc	-0.0853	0.051	-1.686	0.092	-0.184	0.014
term_60 months	0.5553	0.024	23.095	0.000	0.508	0.602
verification_status_1	0.0955	0.021	4.452	0.000	0.053	0.137
initial_list_status_w	-0.1477	0.020	-7.337	0.000	-0.187	-0.108
application_type_JOINT	-0.2201	0.324	-0.679	0.497	-0.855	0.415
grade	0.3743	0.099	3.789	0.000	0.181	0.568
title	6.3559	0.193	32.961	0.000	5.978	6.734
emp_title	6.8480	0.098	70.114	0.000	6.657	7.039
pincode	38.3578	1.448	26.493	0.000	35.520	41.196
emp_length	0.2255	0.041	5.542	0.000	0.146	0.305
home_ownership	0.2336	0.021	11.115	0.000	0.192	0.275
purpose	-0.2294	0.077	-2.982	0.003	-0.380	-0.079
issue_month	0.1419	0.032	4.371	0.000	0.078	0.206
issue_year	0.4513	0.029	15.346	0.000	0.394	0.509
earliest_cr_line_month	0.0619	0.032	1.933	0.053	-0.001	0.125

In [135]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_v
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[135]:

	feature	VIF
0	int_rate	39.48
10	grade	22.39
11	title	13.47
16	purpose	11.85
2	annual_inc	8.15
3	dti	7.35
5	total_acc	7.31
4	revol_util	7.09
1	installment	6.00
17	issue_month	4.95
12	emp_title	4.61
18	issue_year	4.60
19	earliest_cr_line_month	3.39
7	verification_status_1	2.97
15	home_ownership	2.30
8	initial_list_status_w	1.77
6	term_60 months	1.68
13	pincode	1.62
14	emp_length	1.30
9	application_type_JOINT	1.00

In [136]:

```

## P-value and VIF for earliest_cr_line_month is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL', 'earliest_cr_line_year', 'earliest_
y = df['loan_status']

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob=logsk.predict_proba(X_test)

#Probabilities
print(y_pred,y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test,y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[9.31575462e-01 6.84245378e-02]
 [8.46215477e-01 1.53784523e-01]
 [9.99583939e-01 4.16061182e-04]

```

```
...
```

```

[9.98930911e-01 1.06908902e-03]
 [8.91838648e-01 1.08161352e-01]
 [8.39306554e-01 1.60693446e-01]]

```

```
****Score is :0.9290707769683534****
```

	precision	recall	f1-score	support
0.0	0.90	0.99	0.94	58474
1.0	0.90	0.53	0.67	13095
accuracy			0.90	71569
macro avg	0.90	0.76	0.81	71569
weighted avg	0.90	0.90	0.89	71569

In [137]:

```
#Logistic regression model
logm5 = sm.GLM(y_train,(sm.add_constant(X_train)),family=sm.families.Binomial())
logm5.fit().summary()
```

Out[137]:

Generalized Linear Model Regression Results

Dep. Variable:	loan_status	No. Observations:	166992
Model:	GLM	Df Residuals:	166972
Model Family:	Binomial	Df Model:	19
Link Function:	logit	Scale:	1.0000
Method:	IRLS	Log-Likelihood:	-37373.
Date:	Mon, 28 Nov 2022	Deviance:	74747.
Time:	22:30:00	Pearson chi2:	8.93e+04
No. Iterations:	12	Pseudo R-squ. (CS):	0.3925
Covariance Type:	nonrobust		

	coef	std err	z	P> z	[0.025	0.975]
const	-12.9357	0.275	-47.103	0.000	-13.474	-12.397
int_rate	1.4381	0.128	11.270	0.000	1.188	1.688
installment	0.3569	0.049	7.240	0.000	0.260	0.453
annual_inc	-0.7668	0.064	-11.915	0.000	-0.893	-0.641
dti	0.5604	0.046	12.213	0.000	0.470	0.650
revol_util	0.4059	0.042	9.732	0.000	0.324	0.488
total_acc	-0.0872	0.051	-1.724	0.085	-0.186	0.012
term_60 months	0.5551	0.024	23.088	0.000	0.508	0.602
verification_status_1	0.0957	0.021	4.466	0.000	0.054	0.138
initial_list_status_w	-0.1472	0.020	-7.312	0.000	-0.187	-0.108
application_type_JOINT	-0.2177	0.324	-0.672	0.502	-0.853	0.418
grade	0.3754	0.099	3.801	0.000	0.182	0.569
title	6.3576	0.193	32.971	0.000	5.980	6.736
emp_title	6.8479	0.098	70.110	0.000	6.656	7.039
pincode	38.3601	1.448	26.489	0.000	35.522	41.198
emp_length	0.2261	0.041	5.559	0.000	0.146	0.306
home_ownership	0.2332	0.021	11.097	0.000	0.192	0.274
purpose	-0.2292	0.077	-2.980	0.003	-0.380	-0.078
issue_month	0.1417	0.032	4.365	0.000	0.078	0.205
issue_year	0.4524	0.029	15.384	0.000	0.395	0.510

In [138]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_v
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[138]:

	feature	VIF
0	int_rate	39.34
10	grade	22.32
11	title	13.37
16	purpose	11.79
2	annual_inc	8.10
3	dti	7.33
5	total_acc	7.30
4	revol_util	7.08
1	installment	6.00
17	issue_month	4.94
12	emp_title	4.59
18	issue_year	4.59
7	verification_status_1	2.97
15	home_ownership	2.30
8	initial_list_status_w	1.76
6	term_60 months	1.68
13	pincode	1.62
14	emp_length	1.30
9	application_type_JOINT	1.00

In [139]:

```

## P-value and VIF for title is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL', 'earliest_cr_line_year', 'earliest_
y = df['loan_status']

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob=logsk.predict_proba(X_test)

#Probabilities
print(y_pred,y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test,y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[9.13513892e-01 8.64861083e-02]
 [7.93064372e-01 2.06935628e-01]
 [9.99627307e-01 3.72693493e-04]
 ...
 [9.98925582e-01 1.07441793e-03]
 [8.92259668e-01 1.07740332e-01]
 [8.75694580e-01 1.24305420e-01]]

```

```

****Score is :0.9252781944787096****

```

	precision	recall	f1-score	support
0.0	0.90	0.99	0.94	58474
1.0	0.90	0.52	0.66	13095
accuracy			0.90	71569
macro avg	0.90	0.75	0.80	71569
weighted avg	0.90	0.90	0.89	71569

In [140]:

```
#Logistic regression model
logm6 = sm.GLM(y_train,(sm.add_constant(X_train)),family=sm.families.Binomial())
logm6.fit().summary()
```

Out[140]:

Generalized Linear Model Regression Results

Dep. Variable:	loan_status	No. Observations:	166992			
Model:	GLM	Df Residuals:	166973			
Model Family:	Binomial	Df Model:	18			
Link Function:	logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-37985.			
Date:	Mon, 28 Nov 2022	Deviance:	75969.			
Time:	22:30:08	Pearson chi2:	9.00e+04			
No. Iterations:	12	Pseudo R-squ. (CS):	0.3880			
Covariance Type:	nonrobust					
	coef	std err	z	P> z	[0.025	0.975]
const	-12.0456	0.271	-44.434	0.000	-12.577	-11.514
int_rate	1.4247	0.126	11.289	0.000	1.177	1.672
installment	0.3356	0.049	6.857	0.000	0.240	0.432
annual_inc	-0.7578	0.064	-11.854	0.000	-0.883	-0.632
dti	0.5608	0.046	12.302	0.000	0.471	0.650
revol_util	0.3572	0.041	8.634	0.000	0.276	0.438
total_acc	-0.0914	0.050	-1.823	0.068	-0.190	0.007
term_ 60 months	0.5453	0.024	22.856	0.000	0.499	0.592
verification_status_1	0.0991	0.021	4.659	0.000	0.057	0.141
initial_list_status_w	-0.1173	0.020	-5.850	0.000	-0.157	-0.078
application_type_JOINT	-0.1054	0.325	-0.325	0.746	-0.742	0.531
grade	0.4256	0.098	4.356	0.000	0.234	0.617
emp_title	6.8828	0.097	70.961	0.000	6.693	7.073
pincode	38.3091	1.439	26.629	0.000	35.489	41.129
emp_length	0.2340	0.040	5.793	0.000	0.155	0.313
home_ownership	0.2337	0.021	11.212	0.000	0.193	0.275
purpose	0.4979	0.073	6.791	0.000	0.354	0.642
issue_month	0.1448	0.032	4.504	0.000	0.082	0.208
issue_year	0.5787	0.029	19.717	0.000	0.521	0.636

In [141]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_vif1.shape[1])]
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[141]:

	feature	VIF
0	int_rate	39.07
10	grade	22.25
15	purpose	9.89
2	annual_inc	7.95
3	dti	7.29
5	total_acc	7.28
4	revol_util	7.07
1	installment	6.00
16	issue_month	4.87
11	emp_title	4.52
17	issue_year	4.37
7	verification_status_1	2.97
14	home_ownership	2.28
8	initial_list_status_w	1.75
6	term_60 months	1.68
12	pincode	1.62
13	emp_length	1.30
9	application_type_JOINT	1.00

In [142]:

```

## P-value and VIF for total_acc is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL', 'earliest_cr_line_year', 'earliest_
y = df['loan_status']

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob=logsk.predict_proba(X_test)

#Probabilities
print(y_pred,y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test,y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[9.14750597e-01 8.52494030e-02]
 [7.96022222e-01 2.03977778e-01]
 [9.99628399e-01 3.71600557e-04]
 ...
 [9.98909953e-01 1.09004652e-03]
 [8.94148018e-01 1.05851982e-01]
 [8.75346781e-01 1.24653219e-01]]
****Score is :0.9252783250752565****

```

	precision	recall	f1-score	support
0.0	0.90	0.99	0.94	58474
1.0	0.90	0.52	0.66	13095
accuracy			0.90	71569
macro avg	0.90	0.75	0.80	71569
weighted avg	0.90	0.90	0.89	71569

In [143]:

```
#Logistic regression model
logm7 = sm.GLM(y_train,(sm.add_constant(X_train)),family=sm.families.Binomial())
logm7.fit().summary()
```

Out[143]:

Generalized Linear Model Regression Results

Dep. Variable:	loan_status	No. Observations:	166992			
Model:	GLM	Df Residuals:	166974			
Model Family:	Binomial	Df Model:	17			
Link Function:	logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-37986.			
Date:	Mon, 28 Nov 2022	Deviance:	75972.			
Time:	22:30:15	Pearson chi2:	9.00e+04			
No. Iterations:	12	Pseudo R-squ. (CS):	0.3880			
Covariance Type:	nonrobust					
	coef	std err	z	P> z	[0.025	0.975]
const	-12.0691	0.271	-44.562	0.000	-12.600	-11.538
int_rate	1.4225	0.126	11.272	0.000	1.175	1.670
installment	0.3298	0.049	6.754	0.000	0.234	0.426
annual_inc	-0.7899	0.061	-12.845	0.000	-0.910	-0.669
dti	0.5334	0.043	12.391	0.000	0.449	0.618
revol_util	0.3685	0.041	9.003	0.000	0.288	0.449
term_ 60 months	0.5422	0.024	22.786	0.000	0.496	0.589
verification_status_1	0.1001	0.021	4.708	0.000	0.058	0.142
initial_list_status_w	-0.1170	0.020	-5.836	0.000	-0.156	-0.078
application_type_JOINT	-0.1022	0.325	-0.315	0.753	-0.739	0.534
grade	0.4325	0.098	4.430	0.000	0.241	0.624
emp_title	6.8858	0.097	70.994	0.000	6.696	7.076
pincode	38.3106	1.439	26.625	0.000	35.490	41.131
emp_length	0.2323	0.040	5.756	0.000	0.153	0.311
home_ownership	0.2398	0.021	11.659	0.000	0.200	0.280
purpose	0.4969	0.073	6.774	0.000	0.353	0.641
issue_month	0.1448	0.032	4.504	0.000	0.082	0.208
issue_year	0.5780	0.029	19.694	0.000	0.520	0.636

In [144]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_v
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[144]:

	feature	VIF
0	int_rate	38.85
9	grade	22.04
14	purpose	9.75
4	revol_util	6.98
2	annual_inc	6.79
3	dti	6.18
1	installment	5.98
15	issue_month	4.85
10	emp_title	4.51
16	issue_year	4.36
6	verification_status_1	2.97
13	home_ownership	2.24
7	initial_list_status_w	1.75
5	term_60 months	1.68
11	pincode	1.62
12	emp_length	1.30
8	application_type_JOINT	1.00

In [145]:

```

## P-value and VIF for application_type_JOINT is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL', 'earliest_cr_line_year', 'earliest_
y = df['loan_status']

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob=logsk.predict_proba(X_test)

#Probabilities
print(y_pred,y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test,y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[9.14777117e-01 8.52228830e-02]
 [7.96106551e-01 2.03893449e-01]
 [9.99628000e-01 3.72000272e-04]
 ...
 [9.98909168e-01 1.09083163e-03]
 [8.94200763e-01 1.05799237e-01]
 [8.75258940e-01 1.24741060e-01]]
****Score is :0.9252791177962961****

```

	precision	recall	f1-score	support
0.0	0.90	0.99	0.94	58474
1.0	0.90	0.52	0.66	13095
accuracy			0.90	71569
macro avg	0.90	0.75	0.80	71569
weighted avg	0.90	0.90	0.89	71569

In [146]:

```
#Logistic regression model
logm8 = sm.GLM(y_train,(sm.add_constant(X_train)),family=sm.families.Binomial())
logm8.fit().summary()
```

Out[146]:

Generalized Linear Model Regression Results

Dep. Variable:	loan_status	No. Observations:	166992			
Model:	GLM	Df Residuals:	166975			
Model Family:	Binomial	Df Model:	16			
Link Function:	logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-37986.			
Date:	Mon, 28 Nov 2022	Deviance:	75973.			
Time:	22:30:22	Pearson chi2:	9.00e+04			
No. Iterations:	12	Pseudo R-squ. (CS):	0.3880			
Covariance Type:	nonrobust					
	coef	std err	z	P> z	[0.025	0.975]
const	-12.0692	0.271	-44.563	0.000	-12.600	-11.538
int_rate	1.4231	0.126	11.277	0.000	1.176	1.670
installment	0.3296	0.049	6.750	0.000	0.234	0.425
annual_inc	-0.7897	0.061	-12.843	0.000	-0.910	-0.669
dti	0.5332	0.043	12.388	0.000	0.449	0.618
revol_util	0.3685	0.041	9.004	0.000	0.288	0.449
term_ 60 months	0.5421	0.024	22.784	0.000	0.495	0.589
verification_status_1	0.1001	0.021	4.708	0.000	0.058	0.142
initial_list_status_w	-0.1172	0.020	-5.850	0.000	-0.156	-0.078
grade	0.4320	0.098	4.425	0.000	0.241	0.623
emp_title	6.8857	0.097	70.994	0.000	6.696	7.076
pincode	38.3100	1.439	26.625	0.000	35.490	41.130
emp_length	0.2322	0.040	5.754	0.000	0.153	0.311
home_ownership	0.2399	0.021	11.665	0.000	0.200	0.280
purpose	0.4967	0.073	6.773	0.000	0.353	0.640
issue_month	0.1448	0.032	4.504	0.000	0.082	0.208
issue_year	0.5785	0.029	19.737	0.000	0.521	0.636

In [147]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_vif1.shape[1])]
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[147]:

	feature	VIF
0	int_rate	38.84
8	grade	22.04
13	purpose	9.75
4	revol_util	6.98
2	annual_inc	6.79
3	dti	6.18
1	installment	5.98
14	issue_month	4.85
9	emp_title	4.51
15	issue_year	4.35
6	verification_status_1	2.97
12	home_ownership	2.24
7	initial_list_status_w	1.75
5	term_60 months	1.68
10	pincode	1.62
11	emp_length	1.30

In [148]:

```

## VIF for grade is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL', 'earliest_cr_line_year', 'grade', 'e
y = df['loan_status']

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob=logsk.predict_proba(X_test)

#Probabilities
print(y_pred,y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test,y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[9.14873301e-01 8.51266994e-02]
 [7.91634859e-01 2.08365141e-01]
 [9.99620142e-01 3.79857565e-04]
 ...
 [9.98895473e-01 1.10452701e-03]
 [8.94770140e-01 1.05229860e-01]
 [8.68399493e-01 1.31600507e-01]]
****Score is :0.9251613823973591****

```

	precision	recall	f1-score	support
0.0	0.90	0.99	0.94	58474
1.0	0.90	0.52	0.66	13095
accuracy			0.90	71569
macro avg	0.90	0.75	0.80	71569
weighted avg	0.90	0.90	0.89	71569

In [149]:

```
#Logistic regression model
logm9 = sm.GLM(y_train,(sm.add_constant(X_train)),family=sm.families.Binomial())
logm9.fit().summary()
```

Out[149]:

Generalized Linear Model Regression Results

Dep. Variable:	loan_status	No. Observations:	166992			
Model:	GLM	Df Residuals:	166976			
Model Family:	Binomial	Df Model:	15			
Link Function:	logit	Scale:	1.0000			
Method:	IRLS	Log-Likelihood:	-37996.			
Date:	Mon, 28 Nov 2022	Deviance:	75992.			
Time:	22:30:28	Pearson chi2:	9.00e+04			
No. Iterations:	12	Pseudo R-squ. (CS):	0.3879			
Covariance Type:	nonrobust					
	coef	std err	z	P> z 	[0.025	0.975]
const	-12.1560	0.270	-45.020	0.000	-12.685	-11.627
int_rate	1.9378	0.049	39.378	0.000	1.841	2.034
installment	0.3214	0.049	6.585	0.000	0.226	0.417
annual_inc	-0.7909	0.061	-12.862	0.000	-0.911	-0.670
dti	0.5288	0.043	12.290	0.000	0.444	0.613
revol_util	0.3586	0.041	8.773	0.000	0.279	0.439
term_ 60 months	0.5565	0.024	23.612	0.000	0.510	0.603
verification_status_1	0.1008	0.021	4.738	0.000	0.059	0.142
initial_list_status_w	-0.1134	0.020	-5.667	0.000	-0.153	-0.074
emp_title	6.8913	0.097	71.046	0.000	6.701	7.081
pincode	38.3027	1.438	26.636	0.000	35.484	41.121
emp_length	0.2329	0.040	5.773	0.000	0.154	0.312
home_ownership	0.2420	0.021	11.774	0.000	0.202	0.282
purpose	0.5070	0.073	6.910	0.000	0.363	0.651
issue_month	0.1439	0.032	4.478	0.000	0.081	0.207
issue_year	0.6144	0.028	21.727	0.000	0.559	0.670

In [150]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_v
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[150]:

	feature	VIF
12	purpose	9.62
4	revol_util	6.85
0	int_rate	6.73
2	annual_inc	6.67
3	dti	6.10
1	installment	5.95
13	issue_month	4.82
8	emp_title	4.50
14	issue_year	4.22
6	verification_status_1	2.97
11	home_ownership	2.24
7	initial_list_status_w	1.75
5	term_60 months	1.63
9	pincode	1.62
10	emp_length	1.30

In [151]:

```

## VIF for grade is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL', 'earliest_cr_line_year', 'emp_length'])
y = df['loan_status']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob = logsk.predict_proba(X_test)

#Probabilities
print(y_pred, y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test, y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[9.18635003e-01 8.13649968e-02]
 [8.21304854e-01 1.78695146e-01]
 [9.99621103e-01 3.78896834e-04]
 ...
 [9.98879590e-01 1.12041038e-03]
 [8.92634659e-01 1.07365341e-01]
 [8.67197915e-01 1.32802085e-01]]
****Score is :0.9251169142731486****

```

	precision	recall	f1-score	support
0.0	0.90	0.99	0.94	58474
1.0	0.91	0.52	0.66	13095
accuracy			0.90	71569
macro avg	0.90	0.75	0.80	71569
weighted avg	0.90	0.90	0.89	71569

In [152]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_v
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[152]:

	feature	VIF
11	purpose	9.60
4	revol_util	6.85
0	int_rate	6.73
2	annual_inc	6.63
3	dti	6.10
1	installment	5.95
12	issue_month	4.82
8	emp_title	4.42
13	issue_year	4.22
6	verification_status_1	2.90
10	home_ownership	2.24
7	initial_list_status_w	1.75
5	term_60 months	1.63
9	pincode	1.62

In [153]:

```

## VIF for grade is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL', 'earliest_cr_line_year', 'emp_title'])
y = df['loan_status']

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob=logsk.predict_proba(X_test)

#Probabilities
print(y_pred,y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test,y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[9.25482412e-01 7.45175881e-02]
 [8.58210822e-01 1.41789178e-01]
 [9.98613013e-01 1.38698728e-03]
 ...
 [9.99396176e-01 6.03824117e-04]
 [9.21794278e-01 7.82057217e-02]
 [8.45114527e-01 1.54885473e-01]]
****Score is :0.9055412558866556****

```

	precision	recall	f1-score	support
0.0	0.89	0.99	0.94	58474
1.0	0.95	0.46	0.62	13095
accuracy			0.90	71569
macro avg	0.92	0.72	0.78	71569
weighted avg	0.90	0.90	0.88	71569

In [154]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_v
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[154]:

	feature	VIF
10	purpose	9.35
4	revol_util	6.84
0	int_rate	6.68
2	annual_inc	6.63
3	dti	6.03
1	installment	5.95
11	issue_month	4.78
12	issue_year	4.15
6	verification_status_1	2.90
9	home_ownership	2.23
7	initial_list_status_w	1.74
5	term_60 months	1.62
8	pincode	1.57

In [155]:

```

## VIF for purpose is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL', 'earliest_cr_line_year', 'purpose'],
y = df['loan_status']

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob=logsk.predict_proba(X_test)

#Probabilities
print(y_pred,y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test,y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[9.27755222e-01 7.22447777e-02]
 [8.62912095e-01 1.37087905e-01]
 [9.98659638e-01 1.34036194e-03]
 ...
 [9.99352514e-01 6.47485594e-04]
 [9.24423709e-01 7.55762910e-02]
 [8.35926606e-01 1.64073394e-01]]

```

```

****Score is :0.9052690548099736****

```

	precision	recall	f1-score	support
0.0	0.89	0.99	0.94	58474
1.0	0.95	0.46	0.62	13095
accuracy			0.90	71569
macro avg	0.92	0.73	0.78	71569
weighted avg	0.90	0.90	0.88	71569

In [156]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_v
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[156]:

	feature	VIF
4	revol_util	6.82
0	int_rate	6.29
2	annual_inc	6.05
1	installment	5.93
3	dti	5.73
10	issue_month	4.62
11	issue_year	4.10
6	verification_status_1	2.88
9	home_ownership	2.16
7	initial_list_status_w	1.74
5	term_60 months	1.61
8	pincode	1.57

In [157]:

```

## VIF for annual_inc is high, removing the column

X = df.drop(['loan_status', 'application_type_INDIVIDUAL', 'earliest_cr_line_year', 'purpose'],
y = df['loan_status']

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.3,random_state=100)

logsk = LogisticRegression()
logsk.fit(X_train, y_train)

y_pred = logsk.predict(X_test)
y_pred_prob=logsk.predict_proba(X_test)

#Probabilities
print(y_pred,y_pred_prob)

#score
print('****Score is :'+ str(roc_auc_score(y_test, y_pred_prob[:,1]))+'****')

print(classification_report(y_test,y_pred))

```

```

[0. 0. 0. ... 0. 0. 0.] [[0.90966888 0.09033112]
 [0.87245463 0.12754537]
 [0.99869482 0.00130518]
 ...
 [0.99889543 0.00110457]
 [0.89346359 0.10653641]
 [0.86144961 0.13855039]]
****Score is :0.9033841888040546****

```

	precision	recall	f1-score	support
0.0	0.89	1.00	0.94	58474
1.0	0.96	0.45	0.61	13095
accuracy			0.90	71569
macro avg	0.92	0.72	0.78	71569
weighted avg	0.90	0.90	0.88	71569

In [158]:

```
xtrain_vif1 = X_train
vif_data = pd.DataFrame()
vif_data["feature"] = xtrain_vif1.columns
vif_data["VIF"] = [variance_inflation_factor(xtrain_vif1.values, i) for i in range(xtrain_v
vif_data['VIF'] = round(vif_data['VIF'],2)
vif_data = vif_data.sort_values(by = 'VIF',ascending=False)
vif_data
```

Out[158]:

	feature	VIF
3	revol_util	6.47
0	int_rate	6.27
2	dti	5.72
1	installment	4.62
9	issue_month	4.39
10	issue_year	4.01
5	verification_status_1	2.87
8	home_ownership	2.16
6	initial_list_status_w	1.73
4	term_60 months	1.60
7	pincode	1.57

In [159]:

logsk.coef_

Out[159]:

```
array([[ 2.19782436, -0.09273059,  0.74539676,  0.21405641,  0.42921008,
         0.15282871, -0.0827971 , 24.64755252,  0.2997662 ,  0.12556166,
         0.77894413]])
```

In [160]:

```
col = ['int_rate', 'installment', 'dti', 'revol_util', 'term_60 months',
       'verification_status_1', 'initial_list_status_w', 'pincode',
       'home_ownership', 'issue_month', 'issue_year']
```


In [161]:

```
for idx, col_name in enumerate(X_train[col].columns):
    print("The coefficient for {} is {}".format(col_name, np.round(logsk.coef_[0][idx],2)))
```

The coefficient for int_rate is 2.2
 The coefficient for installment is -0.09
 The coefficient for dti is 0.75
 The coefficient for revol_util is 0.21
 The coefficient for term_60 months is 0.43
 The coefficient for verification_status_1 is 0.15
 The coefficient for initial_list_status_w is -0.08
 The coefficient for pincode is 24.65
 The coefficient for home_ownership is 0.3
 The coefficient for issue_month is 0.13
 The coefficient for issue_year is 0.78

We can infer that pincode is very important column

Classification Report (Confusion Matrix,ROCAUC Score)

In [162]:

```
metrics.confusion_matrix(y_test, y_pred)
```

Out[162]:

```
array([[58211, 263],
       [ 7201, 5894]], dtype=int64)
```

In [163]:

```
print(classification_report(y_pred,y_test))
```

```
roc_auc_score(y_test, y_pred_prob[:,1])
```

	precision	recall	f1-score	support
0.0	1.00	0.89	0.94	65412
1.0	0.45	0.96	0.61	6157
accuracy			0.90	71569
macro avg	0.72	0.92	0.78	71569
weighted avg	0.95	0.90	0.91	71569

Out[163]:


```
0.9033841888040546
```

In [164]:

```
fpr,tpr,threshold=roc_curve(y_test,y_pred_prob[:,1])
# roc curve for tpr = fpr
random_probs = [0 for i in range(len(y_test))]
p_fpr, p_tpr, _ = roc_curve(y_test, random_probs, pos_label=1)
plt.plot(fpr,tpr)
plt.plot(p_fpr, p_tpr, linestyle=':', color='red')
plt.title('ROC curve')
# x label
plt.xlabel('False Positive Rate')
# y label
plt.ylabel('True Positive rate')
```

Out[164]:

Text(0, 0.5, 'True Positive rate')




The ROC curve is above the AVG line, we can interpret that the model can classify the between the defaulters

In [165]:

```
precision, recall, thresholds = precision_recall_curve(y_test, y_pred_prob[:,1])  
plt.plot(precision, recall)  
plt.xlabel('Recall')  
plt.ylabel('Precision')
```

Out[165]:

Text(0, 0.5, 'Precision')



As the area under the above curve is relatively high, which represents high recall

Tradeoff Questions:

1. How can we make sure that our model can detect real defaulters and there are less false positives?

This is important as we can lose out on an opportunity to finance more individuals and earn interest on it.

From the dataset we can see that the percentage of defaulters is nearly 20%, for which we can perform SMOTE oversampling techniques to get it fixed.

2. Since NPA (non-performing asset) is a real problem in this industry, it's important we play safe and shouldn't disburse loans to anyone.

Loantap should not disburse the loans to every customer, instead we can use the metrics where there will be a sweetspot for precision and recall, by changing the threshold to the point.

Questionnaire

1) What percentage of customers have fully paid their Loan Amount? ---81.84%

2) Comment about the correlation between Loan Amount and Installment features. ----The correlation between them is 0.94 as it is evident that Loan amount is dependent on the Installment amount it depends on the loan amount.

3) The majority of people have home ownership as _____. ----Mortgage

4) People with grades 'A' are more likely to fully pay their loan. (T/F) ----True, it is 93%

5) Name the top 2 afforded job titles. ----Teacher and Manager

6) Thinking from a bank's perspective, which metric should our primary focus be on..ROC,AUC,Precision,Recall,F1 Score ---The F1 Score anyways makes the best choice as it is the harmonic mean of Recall and Precision, and among recall and precision the recall should be high inorder to reduce the false negative and end up losing opportunity.

7) How does the gap in precision and recall affect the bank?
If the Recall value is low ie. False Negative is high, which leads Bank is losing opportunity cost.
If Precision value is low ie. False Positive is high, which leads in increase of defaulters

8) Which were the features that heavily affected the outcome?
----1) What percentage of customers have fully paid their Loan Amount? ---81.84%

2) Comment about the correlation between Loan Amount and Installment features. ----The correlation between them is 0.94 as it is evident that Loan amount is dependent on the Installment amount it depends on the loan amount.

3) The majority of people have home ownership as _____. ----Mortgage

4) People with grades 'A' are more likely to fully pay their loan. (T/F) ----True, it is 93%

5) Name the top 2 afforded job titles. ----Teacher and Manager

6) Thinking from a bank's perspective, which metric should our primary focus be on..ROC,AUC,Precision,Recall,F1 Score ---The F1 Score anyways makes the best choice as it is the harmonic mean of Recall and Precision, and among recall and precision the recall should be high inorder to reduce the false negative and end up losing opportunity.

7) How does the gap in precision and recall affect the bank?
If the Recall value is low ie. False Negative is high, which leads Bank is losing opportunity cost.
If Precision value is low ie. False Positive is high, which leads in increase of defaulters

8) Which were the features that heavily affected the outcome?
----'int_rate', 'dti', 'pincode', 'issue_year' are the features which are highly affected the outcome.

9) Will the results be affected by geographical location? (Yes/No) ---Yes,Pincode have the greater coefficient which increases the feature importance.

Insights

--The percentage of customers who have fully paid their loan amount is 81.84%

--There is a high correlation between Loan Amount and the Installment

--The majority of people have home ownership as Mortgage

--Grade A people are more likely to pay back the loan when compared to the other grade people

--Teacher and Manager are the top 2 afforded job titles who takes the loan

--'int_rate', 'dti', 'pincode', 'issue_year' are the most significant features

--Higher the interest rate, there is high changes of defaulters as it also has negative coefficients.

--Pincode has higher feature importance compared to other features, loans can be given based on the pincodes too.

Recommendations

In []:

```
--As LoanTap has high percentage of defaulters than other banks, there is high risk of NPAs  
hence LoanTap should try to lower the amount of defaulters to 5%  
--LoanTap should provide slightly higher interest rate which can help them in offsetting th  
of defaulters and maintain the profitability.  
--Using the model, LoanTap can easily reduce the number of defaulters in their portfolio.  
--As pincode place a very important role, Pincode based market segmentation should be  
included in strategic level  
--As higher loan terms has higher changes of defaulters, LoanTap should focus more on Loan  
for shorter duration ie. 36 months
```